

EDA

Sean

2026-02-23

Context

Train.csv: contains the results of 940 regular season games. Games were played at the venue of the home team; thus, the home team may have had a home team advantage. A description of the game results data table columns, train.csv, is as follows:

GameID = Unique game ID Date = Date of match HomeConf = Home team conference HomeID = Home team ID HomeTeam = Home team name HomePts = Home team pts, e.g., points scored by the home team AwayConf = Away team conference AwayID = Away team ID AwayTeam = Away team name AwayPts = Away team pts, e.g., points scored by the away team HomeWinMargin = Home team points minus away team points, e.g., HomeWinMargin>0 indicates home team won the match, HomeWinMargin<0 indicates away team won the match, HomeWinMargin=0 indicates the game ended in a tie. Predictions.csv: contains seventy-five (75) end-of-season rivalry derby games that are to be played. Contestants must predict the results of these games in the field named “Team1_WinMargin.” For example, Team1_WinMargin = 15 indicates that the user is predicting Team 1 to win by 15 points, Team1_WinMargin = -20 indicates that the user is predicting Team 1 to lose by 20 points (conversely, Team2 is predicted to win by 20 points), and Team1_WinMargin = 0 indicates the game is predicted to end in a tie. Team1_WinMargin must be entered from the perspective of Team 1.

All derby matches are played at a neutral venue with fan support split 50% for team 1 and 50% for team 2; thus, neither team has a home team advantage. A description of the prediction data table data, predictions.csv, is as follows:

GameID = Unique game ID Date = Date of match (to be played) Team1_Conf = Team 1 conference Team1_ID = Team 1 ID Team1 = Team 1 name Team2_Conf = Team 2 conference Team2_ID = Team 2 ID Team2 = Team 2 name Team1_WinMargin = User Predicted Number of Points Team 1 is expected to win. E.g., Team1_WinMargin > 0 indicates Team 1 is predicted to win the game by the specified number of points, Team1_WinMargin < 0 indicates that Team 1 is predicted to lose the game by the specified number of points, Team1_WinMargin = 0 indicates the game is expected to end in a tie.

About: There is a new sport in town - Rocketball Premier League (RPL). This is a combination of soccer, basketball, team handball, football, and cricket. The inaugural season was played in 2025 and included 165 teams across 13 conferences. There were 940 regular season games played from 1/1/25 through 6/30/25. All regular season games were played at the home team’s venue.

At the end of the inaugural season, instead of a traditional US playoff system, RPL has seventy-five (75) European style rivalry “derby” games played at a neutral site. Winning teams earn bragging rights. All derby matches are played on 7/4/25. Only 150 out of the 165 teams chose to participate in the end-of-season derby.

You are hired by the RPL to be a member of their sports analytics team. You are asked to develop a prediction model that can be used to rank teams and predict the winning team and victory margin for the derby matches. Your tasks as an RPL analyst are as follows:

Predict the winning team and victory margin for seventy-five (75) end-of-season rivalry “derby” matches.

Rank all 165 teams using the results of the 940 regular season matches. For example, you must rank teams

from No. 1 (Best) to No. 165 (Worst). The No. 1 (Best) ranked team will win the Rocketball Supporters Shield and be deemed RPL Champion.

Competition: Contestants are provided with a training data set (Train.csv) that includes the results of the 940 season games. Each team played between ten and fourteen regular season games. Games were played against teams in the same conference opponents and against teams in different conference.

Each contestant can develop their own ranking and prediction model. For example, these models can be based on win/loss, margin of victory, points scored, or any combination of these items. Your prediction model can incorporate the team's strength of schedule, results of common opponents, the team's conference strength, etc. All data is provided in the competition datasets.

Common types of sports analytics predictions models include counting, probability and statistical, linear and non-linear regression, weighted and moving averages, probit and logit, fuzzy logic, random forests, neural networks, machine learning, deep learning, and models of models that incorporate the results of different approaches as inputs.

Users can develop a prediction model using a tool of your choice. For example, Python, MATLAB, Excel, VBA, Java, C++.

Mission: Using the game results of the 940 matches (Train.csv file), user will perform the following calculations as part of the competition:

Rank all 165 teams from No. 1 (Best) to No. 165 (Worst). Predict the Win Margin for the 75 rivalry derby matches.

Load Libraries & Data

```
library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.4.3
## Warning: package 'ggplot2' was built under R version 4.4.2
## Warning: package 'readr' was built under R version 4.4.3
## Warning: package 'purrr' was built under R version 4.4.2
## Warning: package 'dplyr' was built under R version 4.4.2
## Warning: package 'forcats' was built under R version 4.4.3
## Warning: package 'lubridate' was built under R version 4.4.2

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4    v readr      2.1.5
## v forcats   1.0.0    v stringr   1.5.1
## v ggplot2   3.5.1    v tibble    3.2.1
## v lubridate 1.9.4    v tidyr    1.3.1
## v purrr    1.0.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(readr)
library(readxl)

## Warning: package 'readxl' was built under R version 4.4.2
```

```

predictions <- read_csv("Predictions.csv")

## Rows: 75 Columns: 9
## -- Column specification -----
## Delimiter: ","
## chr (5): Date, Team1_Conf, Team1, Team2_Conf, Team2
## dbl (3): GameID, Team1_ID, Team2_ID
## lgl (1): Team1_WinMargin
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

rankings    <- read_excel("Rankings.xlsx")
train       <- read_csv("Train.csv")

## Rows: 940 Columns: 11
## -- Column specification -----
## Delimiter: ","
## chr (5): Date, HomeConf, HomeTeam, AwayConf, AwayTeam
## dbl (6): GameID, HomeID, HomePts, AwayID, AwayPts, HomeWinMargin
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

Note: For ranking system, we can use a combination of ELO rating algorithm, win margins, win/loss record, games played, and etc. - To Start, we can bucket all team into their respective conferences and start every team with a fixed base ELO rating, and then calculate ending ELO as the season has played on. We can further stratify based on conference strength, and then adjust the ELO ratings based on the strength of the opponent and the conference overall

Note: For the forecasting/modeling, the goal is predict margin of victory for 75 upcoming derby matches. There are a limited number of parameters and data points to work with; therefore, complexity of the model will be severely limited. Machine learning models will not likely be applicable here. I am thinking Bayesian regression (Conditioning on ELO ratings POSSIBLY); however, any numerical estimator will work and most should be tested along with ensemble models.

Note: STEP 1: To start, we should begin by trying to find value in the data and understanding the data. We can start by looking at the given datasets and using inference techniques to try to find any trends or patterns in the data. Or we can make more advanced columns based on aggregates of other columns.

Note: STEP 2: Next, we will build out the ELO rating algorithm for this fictional league.

Note: STEP 3: Finally, using the ELO ratings, we can build/test multiple regression models to find the best fit for the data.

Note: STEP 4: Then, we optimize our forecasts for the 75 derby matches and submit our predictions.

STEP 1

```

# Explore the data and find any trends or patterns. Begin with inspecting the "train" dataset

train_summary = train %>%
  group_by(HomeTeam) %>%
  summarise(
    GamesPlayed = n(),
    Wins = sum(HomeWinMargin > 0),

```

```

Losses = sum(HomeWinMargin < 0),
Ties = sum(HomeWinMargin == 0),
TotalPointsScored = sum(HomePts),
TotalPointsAllowed = sum(AwayPts),
AverageWinMargin = mean(HomeWinMargin)
) %>%
arrange(desc(Wins))

plot_of_wins = ggplot(train_summary, aes(x = reorder(HomeTeam, Wins), y = Wins)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  coord_flip() +
  labs(title = "Number of Wins by Home Team", x = "Home Team", y = "Number of Wins") +
  theme_minimal()

plot_of_win_margin = ggplot(train_summary, aes(x = reorder(HomeTeam, AverageWinMargin), y = AverageWinMargin)) +
  geom_bar(stat = "identity", fill = "coral") +
  coord_flip() +
  labs(title = "Average Win Margin by Home Team", x = "Home Team", y = "Average Win Margin") +
  theme_minimal()

plot_of_points_scored = ggplot(train_summary, aes(x = reorder(HomeTeam, TotalPointsScored), y = TotalPointsScored)) +
  geom_bar(stat = "identity", fill = "darkgreen") +
  coord_flip() +
  labs(title = "Total Points Scored by Home Team", x = "Home Team", y = "Total Points Scored") +
  theme_minimal()

plot_of_points_allowed = ggplot(train_summary, aes(x = reorder(HomeTeam, TotalPointsAllowed), y = TotalPointsAllowed)) +
  geom_bar(stat = "identity", fill = "darkred") +
  coord_flip() +
  labs(title = "Total Points Allowed by Home Team", x = "Home Team", y = "Total Points Allowed") +
  theme_minimal()

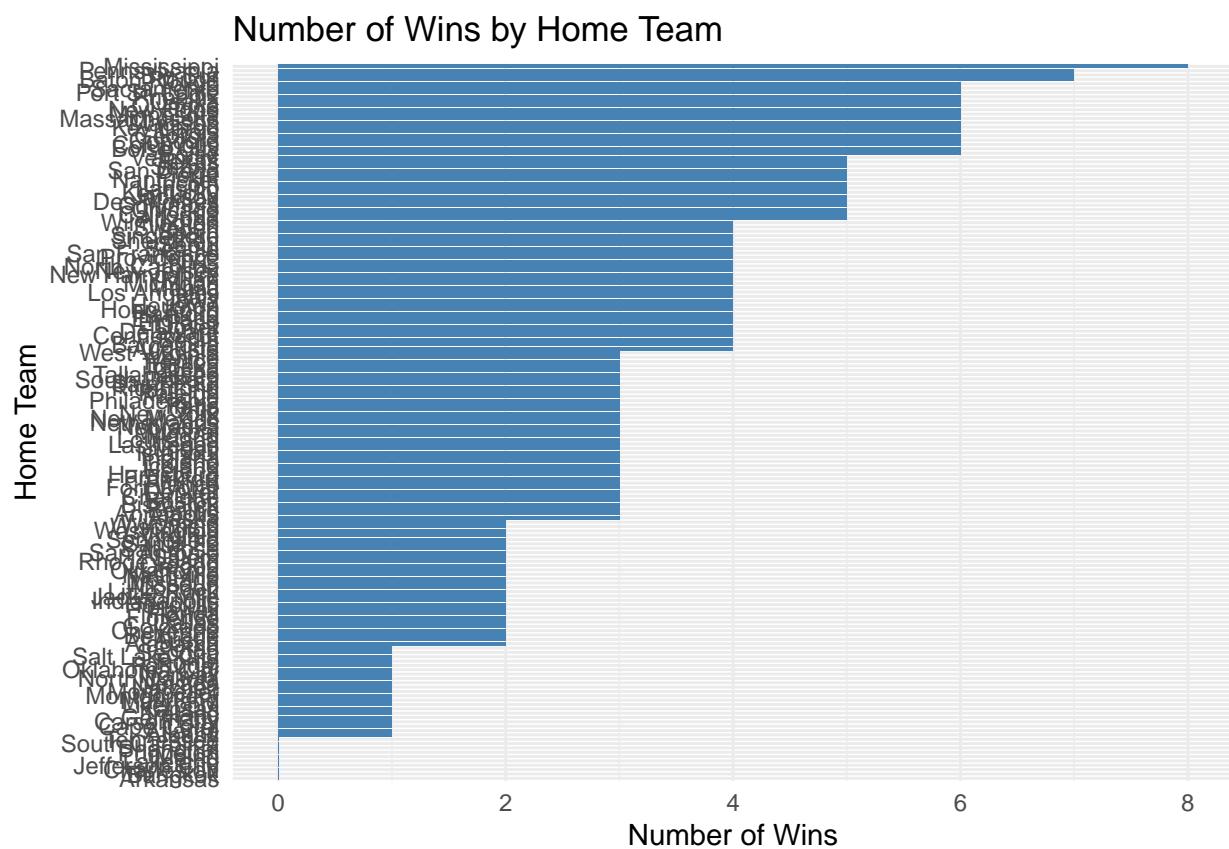
plot_of_games_played = ggplot(train_summary, aes(x = reorder(HomeTeam, GamesPlayed), y = GamesPlayed)) +
  geom_bar(stat = "identity", fill = "purple") +
  coord_flip() +
  labs(title = "Number of Games Played by Home Team", x = "Home Team", y = "Number of Games Played") +
  theme_minimal()

plot_of_conference = ggplot(train, aes(x = HomeConf)) +
  geom_bar(fill = "orange") +
  labs(title = "Number of Games by Conference", x = "Conference", y = "Number of Games") +
  theme_minimal()

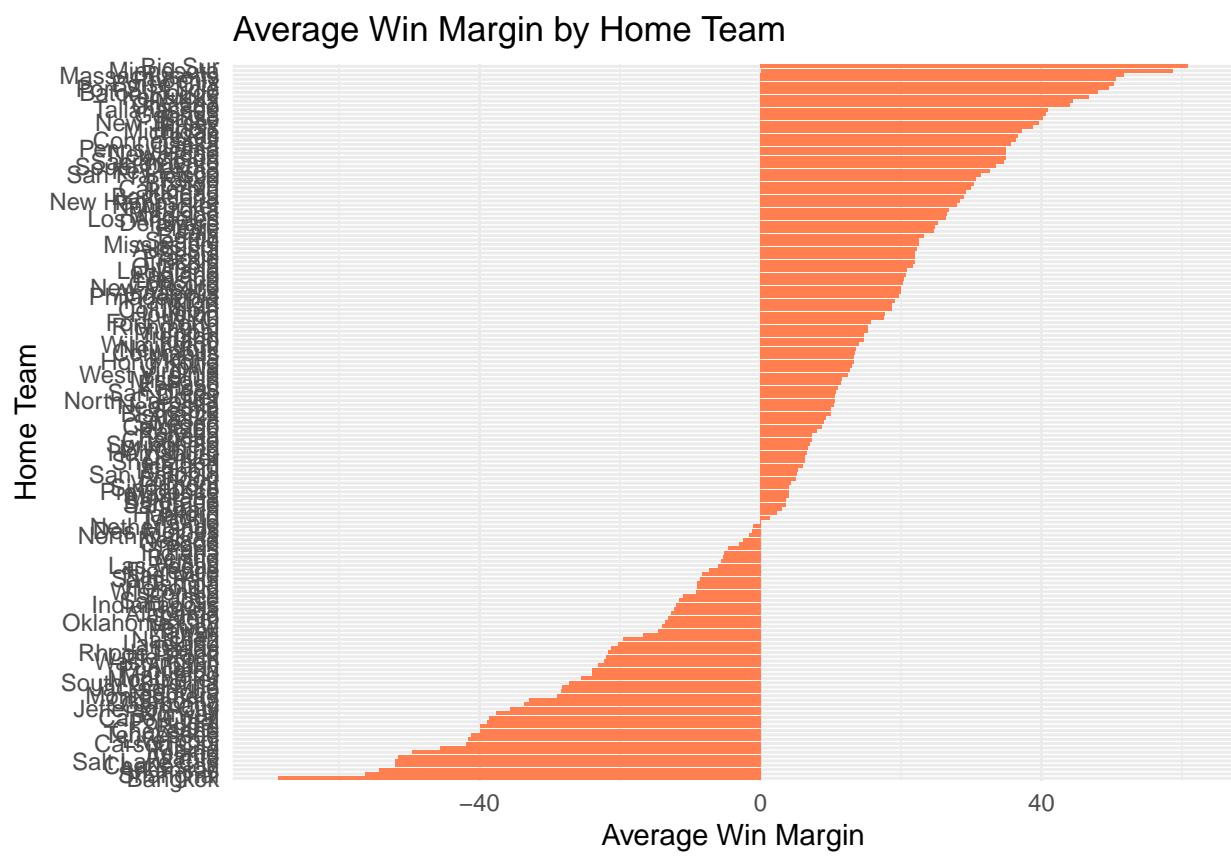
# Display all plots

plot_of_wins

```

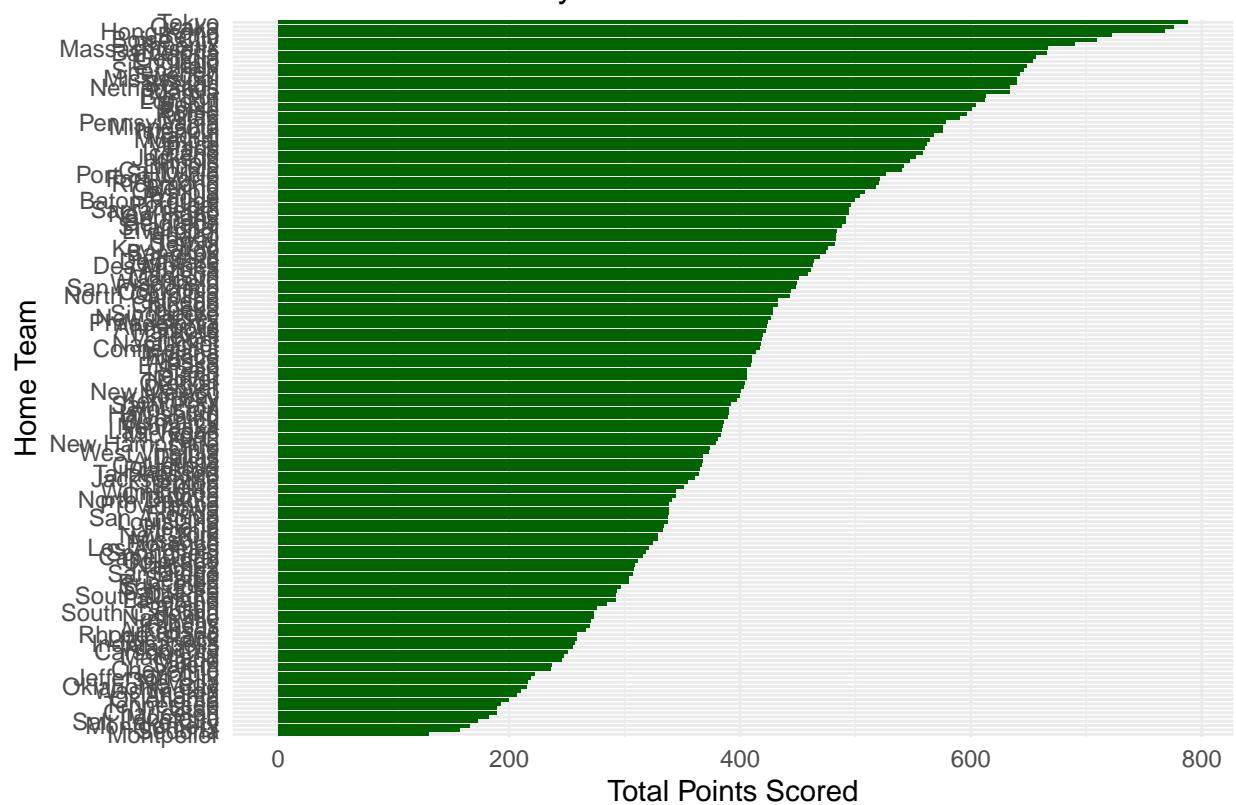


plot_of_win_margin

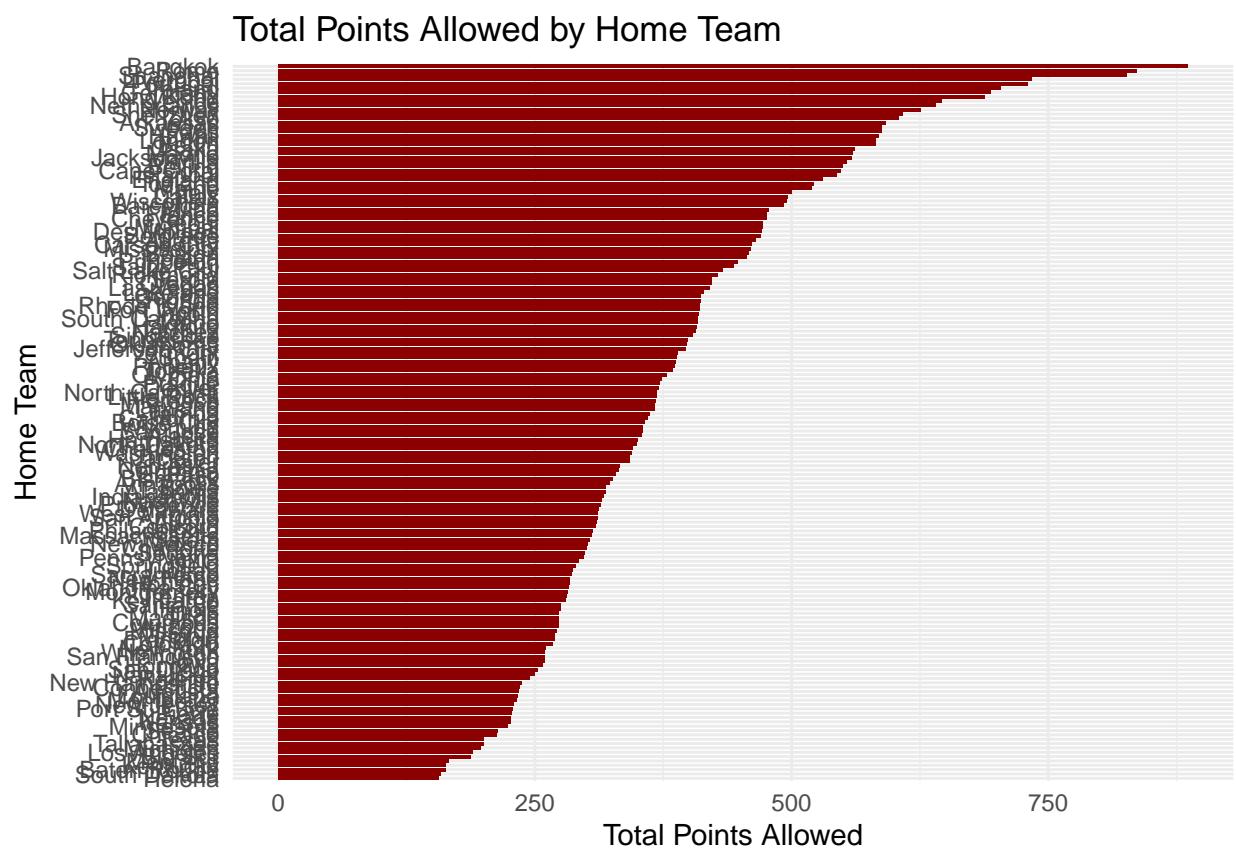


plot_of_points_scored

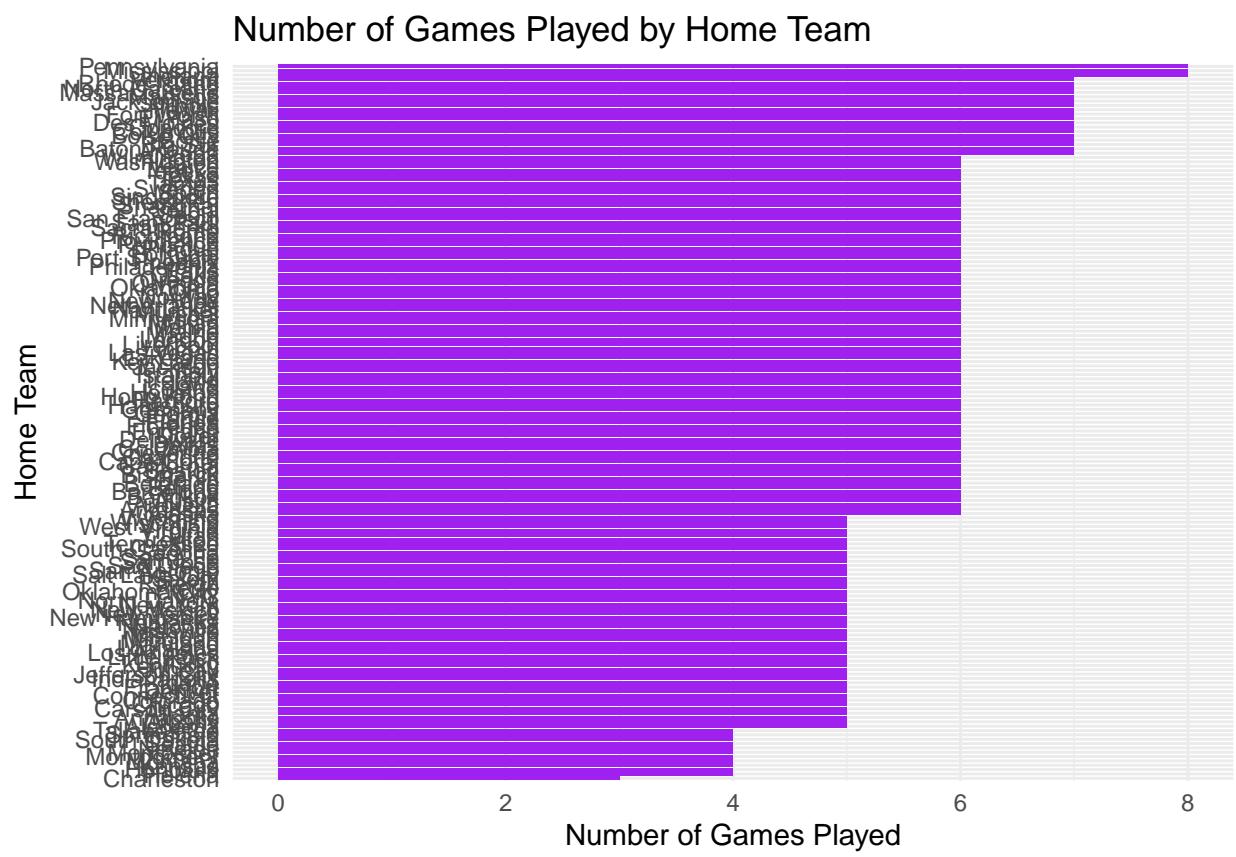
Total Points Scored by Home Team



plot_of_points_allowed

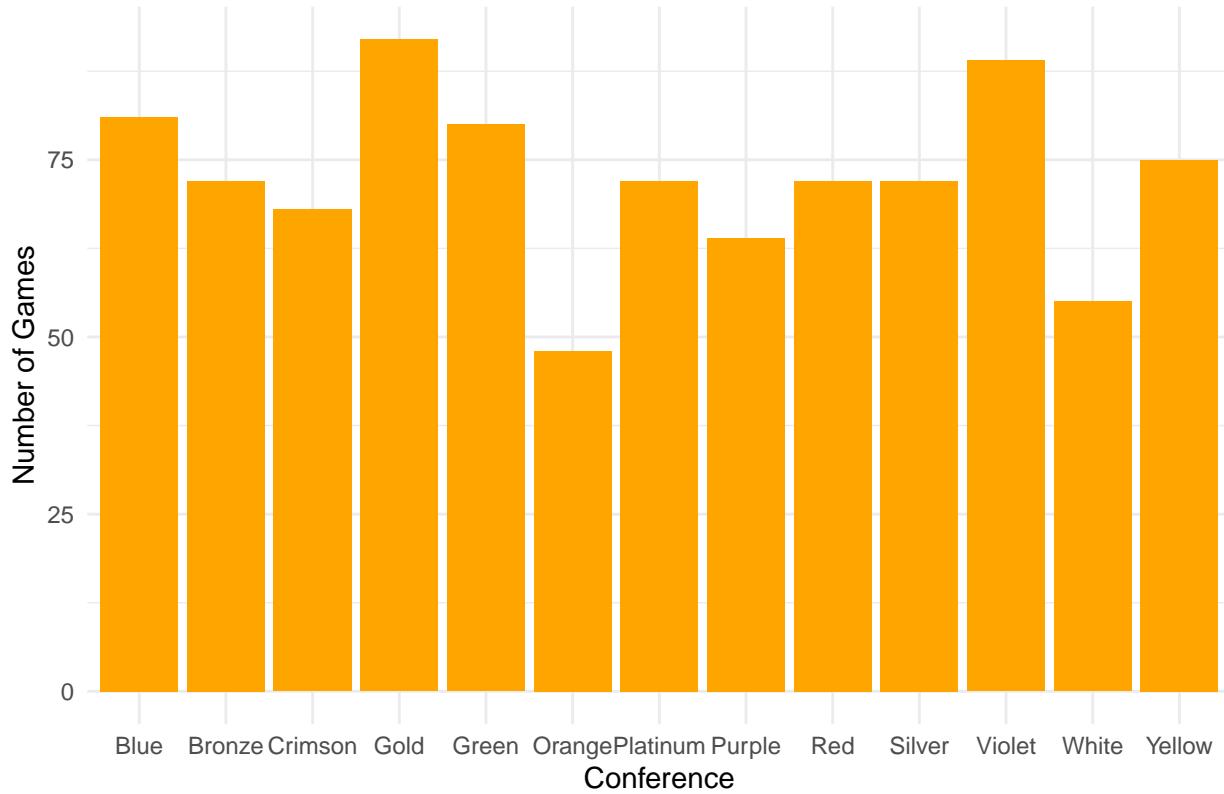


plot_of_games_played



plot_of_conference

Number of Games by Conference



```
# Question: What kinds of inference can be gained from looking at these plots? List the plot and the possible inferences for each.
```

```
# Plot of Wins by Home Team:  
# - This plot shows the distribution of wins among the home teams. We can identify which teams have the highest win rates.
```

```
# Plot of Average Win Margin by Home Team:  
# - This plot shows the average margin of victory for each home team. A higher average win margin may indicate stronger teams.
```

```
# Plot of Total Points Scored by Home Team:  
# - This plot shows the total points scored by each home team. It can indicate which teams have strong offensive abilities.
```

```
# Plot of Total Points Allowed by Home Team:  
# - This plot shows the total points allowed by each home team. It can indicate which teams have strong defensive abilities.
```

```
# Plot of Number of Games Played by Home Team:  
# - This plot shows how many games each home team has played. It can indicate if there are any teams that have played significantly more or fewer games.
```

```
# Plot of Number of Games by Conference:  
# - This plot shows the distribution of games across different conferences. It can indicate if there are any conferences that play significantly more or fewer games.
```



```
# Question 1.1: For the "Number of Wins by Home Team" plot, what numerical distinction would you use to separate the top performing teams from the rest?
```

```
# Question 1.2: For the "Average Win Margin by Home Team" plot, what numerical distinction would you use to separate the teams with the largest margins from the rest?
```

```
# Question 1.3: For the "Total Points Scored by Home Team" plot, what numerical distinction would you use to separate the teams with the highest scoring from the rest?
```

```
# Question 1.4: For the "Total Points Allowed by Home Team" plot, what numerical distinction would you use to separate the teams with the lowest scoring from the rest?
```

```
# Question 1.5: For the "Number of Games Played by Home Team" plot, what kinds of advantages/disadvantages might be associated with playing more or fewer games?
```

```

# Question 1.6: For the "Number of Games Played by Conference" plot, what kinds of advantages/disadvantages do you see? Answer in the next section.

# Answers for Questions 1.1 - 1.6 (make sure to include reasons "why/how" for the "numerical distinctions" between the two groups)

# 1.1: To identify "dominant teams" in the "Number of Wins by Home Team" plot, we could set a numerical threshold for the number of wins.

# 1.1 in R:

win_threshold = quantile(train_summary$Wins, 0.75)
dominant_teams = train_summary %>%
  filter(Wins > win_threshold) %>%
  select(HomeTeam, Wins)
dominant_teams

## # A tibble: 36 x 2
##   HomeTeam     Wins
##   <chr>      <int>
## 1 Mississippi    8
## 2 Baton Rouge    7
## 3 Big Sur        7
## 4 Pennsylvania   7
## 5 Boise          6
## 6 Boise City     6
## 7 Columbus        6
## 8 Concord         6
## 9 Georgia         6
## 10 Illinois       6
## # i 26 more rows

# 1.2: To identify "strong teams" vs "weak teams" in the "Average Win Margin by Home Team" plot, we could set a numerical threshold for the average win margin.

# 1.2 in R:

win_margin_threshold = 10
strong_teams = train_summary %>%
  filter(AverageWinMargin > win_margin_threshold) %>%
  select(HomeTeam, AverageWinMargin)
weak_teams = train_summary %>%
  filter(AverageWinMargin <= win_margin_threshold) %>%
  select(HomeTeam, AverageWinMargin)
strong_teams

## # A tibble: 79 x 2
##   HomeTeam     AverageWinMargin
##   <chr>           <dbl>
## 1 Mississippi     22.5
## 2 Baton Rouge      48
## 3 Big Sur        60.9
## 4 Pennsylvania     35
## 5 Boise          40.3
## 6 Boise City      50.3
## 7 Columbus         13.4

```

```

## 8 Concord          20.3
## 9 Georgia          40.7
## 10 Illinois         38.9
## # i 69 more rows
weak_teams

```

```

## # A tibble: 86 x 2
##   HomeTeam   AverageWinMargin
##   <chr>           <dbl>
## 1 Des Moines      -1.14
## 2 Serbia            10
## 3 Vermont           4.29
## 4 Hartford          2.33
## 5 Providence         4
## 6 Shenzhen          6.33
## 7 Singapore          4
## 8 Spain              3
## 9 Sweden              9
## 10 Alaska            9.33
## # i 76 more rows

```

1.3: To identify "strong offenses" in the "Total Points Scored by Home Team" plot, we could set a num

1.3 in R:

```

points_scored_threshold = quantile(train_summary$TotalPointsScored, 0.75)
strong_offenses = train_summary %>%
  filter(TotalPointsScored > points_scored_threshold) %>%
  select(HomeTeam, TotalPointsScored)
strong_offenses

```

```

## # A tibble: 41 x 2
##   HomeTeam   TotalPointsScored
##   <chr>           <dbl>
## 1 Mississippi      640
## 2 Big Sur           613
## 3 Pennsylvania      578
## 4 Boise             601
## 5 Boise City        709
## 6 Georgia            656
## 7 Illinois           547
## 8 Massachusetts      667
## 9 Minnesota          576
## 10 Olympia            508
## # i 31 more rows

```

1.4: To identify "strong defenses" in the "Total Points Allowed by Home Team" plot, we could set a nu

1.4 in R:

```

points_allowed_threshold = quantile(train_summary$TotalPointsAllowed, 0.25)
strong_defenses = train_summary %>%
  filter(TotalPointsAllowed < points_allowed_threshold) %>%
  select(HomeTeam, TotalPointsAllowed)
strong_defenses

```

```

## # A tibble: 41 x 2
##   HomeTeam    TotalPointsAllowed
##   <chr>          <dbl>
## 1 Baton Rouge      163
## 2 Big Sur          187
## 3 Columbus         273
## 4 Illinois         275
## 5 Madison          273
## 6 Minnesota        224
## 7 Port St. Lucie   228
## 8 Arizona          273
## 9 Chicago           213
## 10 Kentucky         163
## # i 31 more rows

# 1.5: Playing more games can have both advantages and disadvantages for a team's performance in future

# 1.5 in R:

train_summary = train_summary %>%
  mutate(FatigueIndex = GamesPlayed / max(GamesPlayed)) # Example of a fatigue index based on the number of games played

# 1.6: Playing more games within a conference can have advantages such as increased familiarity with opponents

# 1.6 in R:

conference_strength = train %>%
  group_by(HomeConf) %>%
  summarise(
    AverageWinMargin = mean(HomeWinMargin),
    TotalPointsScored = sum(HomePts),
    TotalPointsAllowed = sum(AwayPts)
  ) %>%
  arrange(desc(AverageWinMargin))
conference_strength = conference_strength %>%
  mutate(StrengthOfSchedule = AverageWinMargin / max(AverageWinMargin)) # Example of a strength of schedule index

# Question 1.7: What other trends, patterns, inference, or insights can you think of to gain from exploring this data?

# 1.7: Additional questions to explore the data:

# - Is there a correlation between the number of games played and the average win margin for home teams?

# Answer: If there is a positive correlation, it may suggest that teams that play more games have better win margins.

# Explore answer in R:

correlation_games_win_margin = cor(train_summary$GamesPlayed, train_summary$AverageWinMargin)
correlation_games_win_margin

## [1] 0.1564232

```

```

# - How does the performance of teams vary across different conferences? This could help identify if certain conferences consistently have higher average win margins or total points scored, indicating competitive balance.

# Answer: If certain conferences consistently have higher average win margins or total points scored, it suggests competitive balance.

# Explore answer in R:

conference_performance = train %>%
  group_by(HomeConf) %>%
  summarise(
    AverageWinMargin = mean(HomeWinMargin),
    TotalPointsScored = sum(HomePts),
    TotalPointsAllowed = sum(AwayPts)
  ) %>%
  arrange(desc(AverageWinMargin))
conference_performance

## # A tibble: 13 x 4
##   HomeConf AverageWinMargin TotalPointsScored TotalPointsAllowed
##   <chr>          <dbl>            <dbl>            <dbl>
## 1 Red              19.5            5498            4094
## 2 Blue             11.7            5267            4317
## 3 Violet           10.7            5927            4971
## 4 Gold              10.5            5692            4722
## 5 Green             9.48            5772            5014
## 6 Crimson           8.29            5118            4554
## 7 Orange             4.21            3258            3056
## 8 White              2.33            3477            3349
## 9 Platinum           1.31            7264            7170
## 10 Purple             1                3944            3880
## 11 Yellow            0.427            4609            4577
## 12 Bronze            -0.389           6680            6708
## 13 Silver             -8.22            6228            6820

# - Are there any teams that consistently perform well against specific opponents or in certain conference matchups?

# Answer: If certain teams consistently perform well against specific opponents or in certain conference matchups, it suggests competitive balance.

# Explore the answer in R:

win_margin_distribution = train %>%
  group_by(HomeTeam) %>%
  summarise(
    AverageWinMargin = mean(HomeWinMargin),
    WinMarginVariance = var(HomeWinMargin)
  ) %>%
  arrange(desc(AverageWinMargin))
win_margin_distribution

## # A tibble: 165 x 3

```

```

##   HomeTeam      AverageWinMargin WinMarginVariance
##   <chr>          <dbl>            <dbl>
## 1 Big Sur        60.9             1150.
## 2 Minnesota      58.7             897.
## 3 Massachusetts  51.7             2246.
## 4 Phoenix         50.7             1041.
## 5 Boise City     50.3             1305.
## 6 Port St. Lucie 49.7             551.
## 7 Baton Rouge    48               984
## 8 Kentucky        46.8             149.
## 9 Helena          44.5             1980.
## 10 Chicago        44               696
## # i 155 more rows
# - How do the points scored and points allowed by teams correlate with their overall win/loss record

# Answer: If there is a strong correlation between points scored/allowed and win/loss record or ranking

points_correlation = train_summary %>%
  summarise(
    CorrelationPointsScoredWins = cor(TotalPointsScored, Wins),
    CorrelationPointsAllowedWins = cor(TotalPointsAllowed, Wins),
    CorrelationPointsScoredRanking = cor(TotalPointsScored, GamesPlayed), # Assuming ranking is based on wins
    CorrelationPointsAllowedRanking = cor(TotalPointsAllowed, GamesPlayed) # Assuming ranking is based on wins
  )
points_correlation

## # A tibble: 1 x 4
##   CorrelationPointsScoredWins CorrelationPointsAllowedW~1 CorrelationPointsScoredRanking CorrelationPointsAllowedRanking
##   <dbl>                  <dbl>                  <dbl>                  <dbl>
## 1 0.609                 -0.317                 0.549
## # i abbreviated names: 1: CorrelationPointsScoredWins,
## #   2: CorrelationPointsScoredRanking
## # i 1 more variable: CorrelationPointsAllowedRanking <dbl>

```

STEP 2: Build ELO Rating System Algorithm - USE THE “elo” package

```

library(elo)

## Warning: package 'elo' was built under R version 4.4.3
# Step 1: Initialize ratings for all 165 teams

ELO_RATINGS_v1 = data.frame(
  Team = unique(c(train$HomeTeam, train$AwayTeam)),
  Rating = 1500 # Starting ELO rating for all teams
)

# Step 2: Update ELO ratings based on game results

for (i in 1:nrow(train)) {

```

```

home_team = train$HomeTeam[i]
away_team = train$AwayTeam[i]
home_points = train$HomePts[i]
away_points = train$AwayPts[i]

# Get current ratings
home_rating = ELO_RATINGS_v1$Rating[ELO_RATINGS_v1$Team == home_team]
away_rating = ELO_RATINGS_v1$Rating[ELO_RATINGS_v1$Team == away_team]

# Calculate expected scores
expected_home = 1 / (1 + 10^((away_rating - home_rating) / 400))
expected_away = 1 / (1 + 10^((home_rating - away_rating) / 400))

# Determine actual scores
if (home_points > away_points) {
  actual_home = 1
  actual_away = 0
} else if (home_points < away_points) {
  actual_home = 0
  actual_away = 1
} else {
  actual_home = 0.5
  actual_away = 0.5
}

# Update ratings
K = 40 # K-factor for rating updates
new_home_rating = home_rating + K * (actual_home - expected_home)
new_away_rating = away_rating + K * (actual_away - expected_away)

ELO_RATINGS_v1$Rating[ELO_RATINGS_v1$Team == home_team] <- new_home_rating
ELO_RATINGS_v1$Rating[ELO_RATINGS_v1$Team == away_team] <- new_away_rating
}

# Step 3: Consider addendums to the ELO ratings based on factors found from initial data analysis
# Goal: Create a data frame of all questions asked and the numerical result(s) from the analysis.
#           Iterate through the data frame and consider whether or not the results from the questions/analy
#           Decide if complexity in the rating system is rewarded or penalized

# Apply results from analysis (Question 1.1-1.6 & 5 questions for 1.7)

# 1.1 - Win Threshold for Dominant Teams:

# Question: Is the current rating system sufficient in identifying dominant teams based on the results

# 1.1 Answer: The current ELO rating system may not be sufficient in identifying dominant teams based s

# 1.2 - Average Win Margin Threshold for Strong vs Weak Teams:

```

```
# Question: Is the current rating system sufficient in identifying strong vs weak teams based on the average points scored by each team?  
# 1.2 Answer: The current ELO rating system may not be sufficient in identifying strong vs weak teams based on the average points scored by each team.  
  
# 1.3 - Total Points Scored Threshold for Strong Offenses:  
# Question: Is the current rating system sufficient in identifying strong offenses based on total points scored by the offense?  
# 1.3 Answer: The current ELO rating system may not be sufficient in identifying strong offenses based on total points scored by the offense.  
  
# 1.4 - Total Points Allowed by Home Team  
# Question: Is the current rating system sufficient in identifying strong defenses? If not, how can we improve it?  
# 1.4 Answer: The current ELO rating system may not be sufficient in identifying strong defenses based on total points allowed by the defense.  
  
# 1.5 - Fatigue Index - Number of Games Played by Home Team  
# Question: Is the current rating system sufficient in accounting for the advantages/disadvantages of playing multiple games?  
# Answer: The current ELO rating system may not be sufficient in accounting for the advantages/disadvantages of playing multiple games.  
# Outline for ELO algo (including addendums from 1.1-1.5)  
  
# 1. Initialize ratings for all 165 teams with a base rating (e.g., 1500).  
# 2. For each game in the training dataset:  
#     a. Retrieve the current ratings of the home and away teams.  
#     b. Calculate the expected scores for both teams based on their current ratings.  
#     c. Determine the actual scores based on the game outcome (win/loss/tie).  
#     d. Update the ratings of both teams using the ELO formula, incorporating any adjustments based on the outcome.  
# 3. After processing all games, the final ELO ratings will reflect the performance of each team throughout the season.  
  
# Question: How do we statistically validate the any of the adjustments made to the ELO algorithm based on the insights about the game outcomes?  
# 1.1 - To statistically validate the adjustments made to the ELO algorithm based on the insights about the game outcomes.  
# 1.2 - To statistically validate the adjustments made to the ELO algorithm based on the insights about the game outcomes.  
# 1.3 - To statistically validate the adjustments made to the ELO algorithm based on the insights about the game outcomes.  
# 1.4 - To statistically validate the adjustments made to the ELO algorithm based on the insights about the game outcomes.  
# 1.5 - To statistically validate the adjustments made to the ELO algorithm based on the insights about the game outcomes.  
  
# UPDATED ELO ALGORITHM:
```

```

# Validation Pipeline:

# Test 1.1 - 1.5 adjustments to the ELO algorithm by comparing predicted outcomes using adjusted ELO ratings

# Test individually and collectively for all 1.1-1.5 adjustments to see which adjustments have the most impact

# Create the adjustments inside of the main algorithm loop that - when called - will test the algorithm

ELO_ALGO_TEST = function(train_data, adjustments = c("dominant_teams", "strong_offenses", "strong_defense")) {
  # Initialize ELO ratings
  ELO_RATINGS = data.frame(
    Team = unique(c(train_data$HomeTeam, train_data$AwayTeam)),
    Rating = 1500 # Starting ELO rating for all teams
  )

  # Store predictions and actual outcomes for validation
  predictions = data.frame(
    HomeTeam = character(),
    AwayTeam = character(),
    PredictedOutcome = numeric(),
    ActualOutcome = numeric()
  )

  for (i in 1:nrow(train_data)) {
    home_team = train_data$HomeTeam[i]
    away_team = train_data$AwayTeam[i]
    home_points = train_data$HomePts[i]
    away_points = train_data$AwayPts[i]

    # Get current ratings
    home_rating = ELO_RATINGS$Rating[ELO_RATINGS$Team == home_team]
    away_rating = ELO_RATINGS$Rating[ELO_RATINGS$Team == away_team]

    # Calculate expected scores
    expected_home = 1 / (1 + 10^((away_rating - home_rating) / 400))
    expected_away = 1 / (1 + 10^((home_rating - away_rating) / 400))

    # Determine actual scores
    if (home_points > away_points) {
      actual_home = 1
      actual_away = 0
    } else if (home_points < away_points) {
      actual_home = 0
      actual_away = 1
    } else {
      actual_home = 0.5
      actual_away = 0.5
    }
  }
}

```

```

# Store predictions for validation
predictions <- rbind(predictions, data.frame(
  HomeTeam = home_team,
  AwayTeam = away_team,
  PredictedOutcome = expected_home,
  ActualOutcome = actual_home
))

# Update ratings with adjustments based on insights from data analysis
K = 40 # K-factor for rating updates

# Adjustments for dominant teams, strong offenses/defenses, and fatigue index can be incorporated here

dominant_teams_adjustment = if ("dominant_teams" %in% adjustments) {
  # Example adjustment for dominant teams
  if (home_team %in% dominant_teams$HomeTeam) {
    home_rating <- home_rating + 20 # Bonus for dominant teams
  }
  if (away_team %in% dominant_teams$HomeTeam) {
    away_rating <- away_rating + 20 # Bonus for dominant teams
  }
}

strong_offenses_adjustment = if ("strong_offenses" %in% adjustments) {
  # Example adjustment for strong offenses
  if (home_team %in% strong_offenses$HomeTeam) {
    home_rating <- home_rating + 10 # Bonus for strong offenses
  }
  if (away_team %in% strong_offenses$HomeTeam) {
    away_rating <- away_rating + 10 # Bonus for strong offenses
  }
}

strong_defenses_adjustment = if ("strong_defenses" %in% adjustments) {
  # Example adjustment for strong defenses
  if (home_team %in% strong_defenses$HomeTeam) {
    home_rating <- home_rating + 10 # Bonus for strong defenses
  }
  if (away_team %in% strong_defenses$HomeTeam) {
    away_rating <- away_rating + 10 # Bonus for strong defenses
  }
}

fatigue_index_adjustment = if ("fatigue_index" %in% adjustments) {
  # Example adjustment for fatigue index
  home_fatigue_index = train_summary$FatigueIndex[train_summary$HomeTeam == home_team]
  away_fatigue_index = train_summary$FatigueIndex[train_summary$HomeTeam == away_team]

  home_rating <- home_rating - (home_fatigue_index * 20) # Penalty for fatigue
  away_rating <- away_rating - (away_fatigue_index * 20) # Penalty for fatigue
}

```

```

new_home_rating = home_rating + K * (actual_home - expected_home)
new_away_rating = away_rating + K * (actual_away - expected_away)

ELO_RATINGS$Rating[ELO_RATINGS$Team == home_team] <- new_home_rating
ELO_RATINGS$Rating[ELO_RATINGS$Team == away_team] <- new_away_rating
}

# Validate predictions using MAE, RMSE, and hypothesis testing
mae = mean(abs(predictions$PredictedOutcome - predictions$ActualOutcome))
rmse = sqrt(mean((predictions$PredictedOutcome - predictions$ActualOutcome)^2))

# Hypothesis testing can be conducted here to compare the adjusted ELO ratings against the original ELO ratings

return(list(
  ELO_RATINGS = ELO_RATINGS,
  Predictions = predictions,
  MAE = mae,
  RMSE = rmse
  # Include results of hypothesis testing here
))
}

# Execution:

results = ELO_ALGO_TEST(train_data = train, adjustments = c("dominant_teams", "strong_defenses", "strong_offenses"))

# Using "results", cross-validate the adjustments to determine which have the most statistical significance

# Question: After the ELO_ALGO_TEST was run and data was collected, what else is needed to test, validate, and refine the algorithm?

# 1. Analyze the results of the ELO_ALGO_TEST, specifically looking at the MAE and RMSE metrics to evaluate the performance of the algorithm
# 2. Conduct hypothesis testing to determine if the adjustments made to the ELO algorithm (for dominant teams, strong defenses, and strong offenses) are statistically significant
# 3. Based on the results of the validation tests, identify which adjustments have the most statistical significance
# 4. Refine the ELO_ALGO_TEST function to create a finalized ELO_ALGO_FINAL function that incorporates the best adjustments

# Step 1 in R:
results = ELO_ALGO_TEST(train_data = train, adjustments = c("dominant_teams", "strong_defenses", "strong_offenses"))
mae = results$MAE
rmse = results$RMSE
print(paste("MAE:", mae))

## [1] "MAE: 0.426608830308534"
print(paste("RMSE:", rmse))

## [1] "RMSE: 0.451006144571495"

```

```

# MAE: 0.426608830308534
# RMSE: 0.451006144571495

# Step 2 in R:

# Hypothesis testing can be conducted using a paired t-test to compare the predicted outcomes from the adjusted ELO ratings and the original ELO ratings.

# Original ELO Predictions:

original_predictions = data.frame(
  HomeTeam = character(),
  AwayTeam = character(),
  PredictedOutcome = numeric(),
  ActualOutcome = numeric()
)
for (i in 1:nrow(train)) {
  home_team = train$HomeTeam[i]
  away_team = train$AwayTeam[i]
  home_points = train$HomePts[i]
  away_points = train$AwayPts[i]

  # Get current ratings from original ELO ratings
  home_rating = ELO_RATINGS_v1$Rating[ELO_RATINGS_v1$Team == home_team]
  away_rating = ELO_RATINGS_v1$Rating[ELO_RATINGS_v1$Team == away_team]

  # Calculate expected scores using original ELO ratings
  expected_home = 1 / (1 + 10^((away_rating - home_rating) / 400))

  # Determine actual scores
  if (home_points > away_points) {
    actual_home = 1
  } else if (home_points < away_points) {
    actual_home = 0
  } else {
    actual_home = 0.5
  }

  # Store predictions for validation
  original_predictions <- rbind(original_predictions, data.frame(
    HomeTeam = home_team,
    AwayTeam = away_team,
    PredictedOutcome = expected_home,
    ActualOutcome = actual_home
  ))
}

# Now we have both the predictions from the adjusted ELO ratings and the results from the original elo

# Paired t-test to compare the predicted outcomes from the adjusted ELO ratings against the actual outcome

t_test_result = t.test(results$Predictions$PredictedOutcome, original_predictions$PredictedOutcome, paired = TRUE)
print(t_test_result)

```

```

## Paired t-test
## data: results$Predictions$PredictedOutcome and original_predictions$PredictedOutcome
## t = -1.7668, df = 939, p-value = 0.07759
## alternative hypothesis: true mean difference is not equal to 0
## 95 percent confidence interval:
## -0.013941973 0.000731784
## sample estimates:
## mean difference
## -0.006605095

# Results:
# t = -1.7668
# df = 939
# p-value = 0.07759
# Sample estimates: Mean Difference: -0.006605095
# 95% CI: [-0.013941973, 0.000731784]

# Step 3 in R - Which Adjustments had the most statistical significance in improving accuracy?

# To determine which adjustments had the most statistical significance in improving accuracy, we would ...

# Run 1: Only "dominant_teams" adjustment

results_dominant_teams = ELO_ALGO_TEST(train_data = train, adjustments = c("dominant_teams"))
mae_dominant_teams = results_dominant_teams$MAE
rmse_dominant_teams = results_dominant_teams$RMSE
print(paste("MAE with Dominant Teams Adjustment:", mae_dominant_teams))

## [1] "MAE with Dominant Teams Adjustment: 0.438924806542782"
print(paste("RMSE with Dominant Teams Adjustment:", rmse_dominant_teams))

## [1] "RMSE with Dominant Teams Adjustment: 0.457196940180702"

# MAE with Dominant Teams: 0.438
# RMSE with Dominant Teams: 0.457

# Run 2: Only "strong_offenses" adjustment

results_strong_offenses = ELO_ALGO_TEST(train_data = train, adjustments = c("strong_offenses"))
mae_strong_offenses = results_strong_offenses$MAE
rmse_strong_offenses = results_strong_offenses$RMSE
print(paste("MAE with Strong Offenses Adjustment:", mae_strong_offenses))

## [1] "MAE with Strong Offenses Adjustment: 0.454860438932721"
print(paste("RMSE with Strong Offenses Adjustment:", rmse_strong_offenses))

## [1] "RMSE with Strong Offenses Adjustment: 0.466590167987282"

# MAE with strong offenses: 0.455
# RMSE with strong offenses: 0.467

# Run 3: Only "strong_defenses" adjustment

```

```

results_strong_defenses = ELO_ALGO_TEST(train_data = train, adjustments = c("strong_defenses"))
mae_strong_defenses = results_strong_defenses$MAE
rmse_strong_defenses = results_strong_defenses$RMSE
print(paste("MAE with Strong Defenses Adjustment:", mae_strong_defenses))

## [1] "MAE with Strong Defenses Adjustment: 0.456683162034806"
print(paste("RMSE with Strong Defenses Adjustment:", rmse_strong_defenses))

## [1] "RMSE with Strong Defenses Adjustment: 0.4687096591706"
# MAE with strong defenses: 0.457
# RMSE with strong defenses: 0.469

# Run 4: Only "fatigue_index" adjustment

results_fatigue_index = ELO_ALGO_TEST(train_data = train, adjustments = c("fatigue_index"))
mae_fatigue_index = results_fatigue_index$MAE
rmse_fatigue_index = results_fatigue_index$RMSE
print(paste("MAE with Fatigue Index Adjustment:", mae_fatigue_index))

## [1] "MAE with Fatigue Index Adjustment: 0.465619602581491"
print(paste("RMSE with Fatigue Index Adjustment:", rmse_fatigue_index))

## [1] "RMSE with Fatigue Index Adjustment: 0.474598493952068"
# MAE with fatigue index: 0.466
# RMSE with fatigue index: 0.475

# Conclusion for Step 3:
# Based on the MAE and RMSE metrics for each individual adjustment, it appears that the "dominant_teams"

# Step 4 in R - Refine the ELO_ALGO_TEST function to create a finalized ELO_ALGO_FINAL function that in

ELO_ALGO_FINAL = function(train_data) {
  # Initialize ELO ratings
  ELO_RATINGS = data.frame(
    Team = unique(c(train_data$HomeTeam, train_data$AwayTeam)),
    Rating = 1500 # Starting ELO rating for all teams
  )

  for (i in 1:nrow(train_data)) {
    home_team = train_data$HomeTeam[i]
    away_team = train_data$AwayTeam[i]
    home_points = train_data$HomePts[i]
    away_points = train_data$AwayPts[i]

    # Get current ratings
    home_rating = ELO_RATINGS$Rating[ELO_RATINGS$Team == home_team]
    away_rating = ELO_RATINGS$Rating[ELO_RATINGS$Team == away_team]

    # Calculate expected scores
    expected_home = 1 / (1 + 10^((away_rating - home_rating) / 400))
  }
}

```

```

# Determine actual scores
if (home_points > away_points) {
  actual_home = 1
  actual_away = 0
} else if (home_points < away_points) {
  actual_home = 0
  actual_away = 1
} else {
  actual_home = 0.5
  actual_away = 0.5
}

# Update ratings with adjustments based on insights from data analysis (only dominant teams adjustment)
K = 40 # K-factor for rating updates

dominant_teams_adjustment = if (home_team %in% dominant_teams$HomeTeam) {
  home_rating <- home_rating + 20 # Bonus for dominant teams
}
if (away_team %in% dominant_teams$HomeTeam) {
  away_rating <- away_rating + 20 # Bonus for dominant teams
}

new_home_rating = home_rating + K * (actual_home - expected_home)
new_away_rating = away_rating + K * (actual_away - expected_away)

ELO_RATINGS$Rating[ELO_RATINGS$Team == home_team] <- new_home_rating
ELO_RATINGS$Rating[ELO_RATINGS$Team == away_team] <- new_away_rating
}

return(ELO_RATINGS)
}

```

STEP 3: Modeling/Forecasting - Create a statistical model that can fit the results of HomeWinMargin and then use that model to predict the HomeWinMargin for the next 75 matchups

```

# Plan:
# 1. Explore the relationship between HomeWinMargin and other variables in the dataset (e.g., HomePts, ...
# 2. Choose an appropriate statistical modeling technique (e.g., linear regression, random forest, etc.)
# 3. Train the chosen model on the training dataset, using HomeWinMargin as the target variable and the ...
# 4. Evaluate the performance of the model using appropriate metrics (e.g., R-squared, mean absolute error, ...)
# 5. Use the trained model to predict the HomeWinMargin for the next 75 matchups, using the relevant predictors

# 3.1 - What predictors should we use? What is the relationship between HomeWinMargin and other variables?
# 3.1 Answer: To identify potential predictors for the model, we can explore the relationships between ...

# Correlation Coefficients:

correlation_home_win_margin = train %>%

```

```

summarise(
  CorrelationHomePts = cor(HomeWinMargin, HomePts),
  CorrelationAwayPts = cor(HomeWinMargin, AwayPts),
  CorrelationHomeConf = cor(HomeWinMargin, as.numeric(as.factor(HomeConf))),
  CorrelationAwayConf = cor(HomeWinMargin, as.numeric(as.factor(AwayConf)))
)
correlation_home_win_margin

## # A tibble: 1 x 4
##   CorrelationHomePts CorrelationAwayPts CorrelationHomeConf CorrelationAwayConf
##       <dbl>           <dbl>           <dbl>           <dbl>
## 1     0.576          -0.640          -0.0391         0.0939
# CorrelationHomePts: 0.5758479
# CorrelationAwayPts: -0.6403852
# CorrelationHomeConf: -0.03906505
# CorrelationAwayConf: 0.09389065

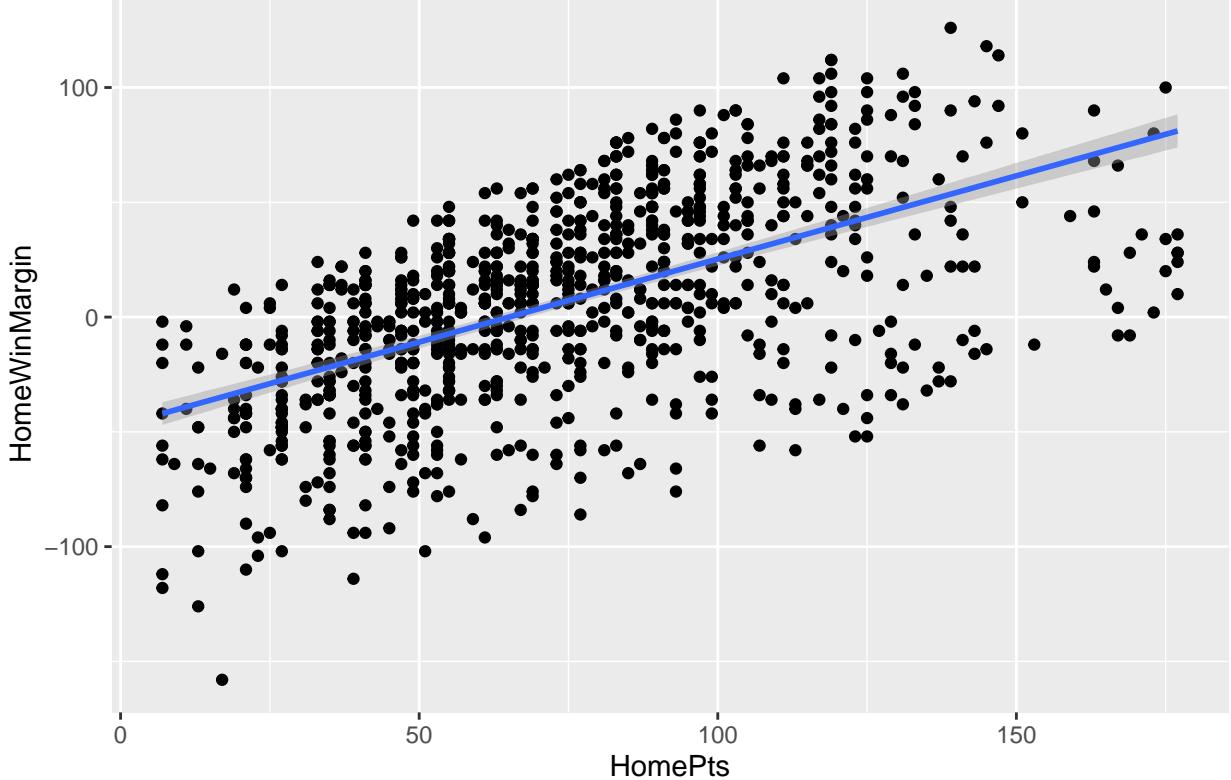
# Scatter & Box Plots:

ggplot(train, aes(x = HomePts, y = HomeWinMargin)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(title = "Scatter Plot of HomePts vs HomeWinMargin")

```

`geom_smooth()` using formula = 'y ~ x'

Scatter Plot of HomePts vs HomeWinMargin

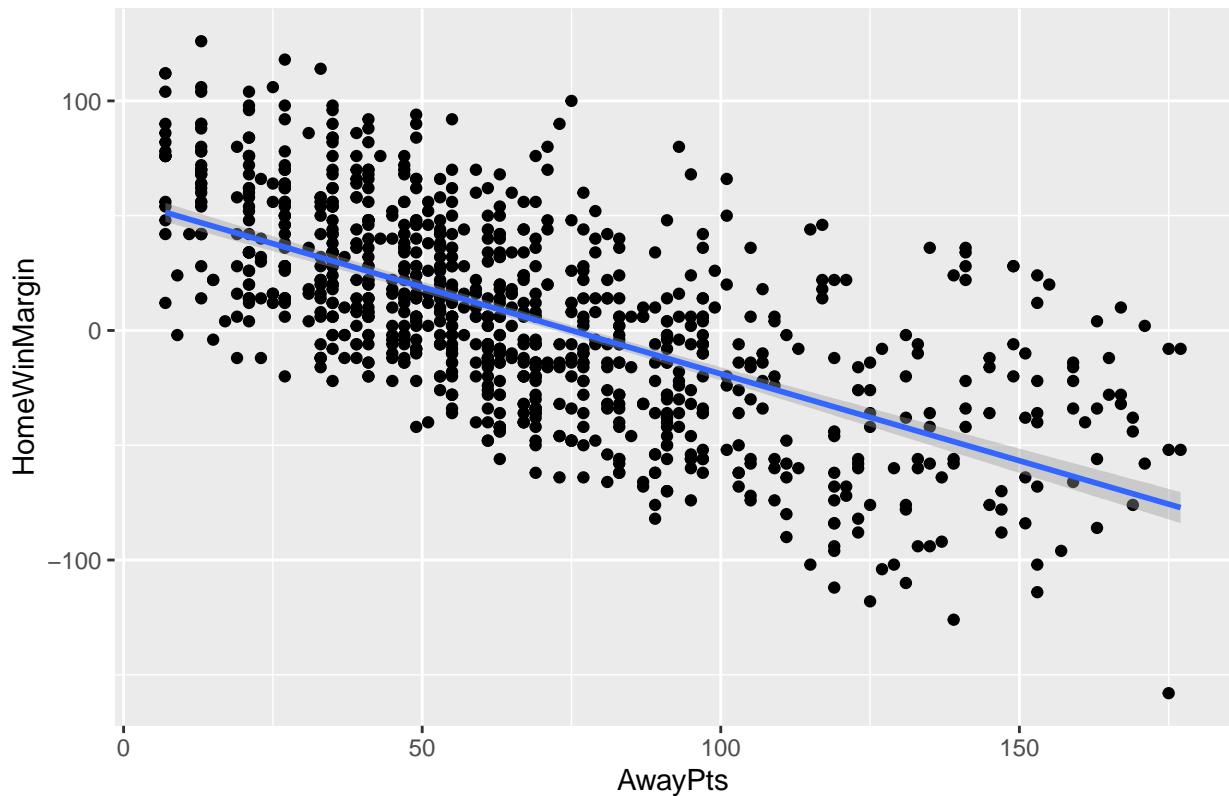


```

ggplot(train, aes(x = AwayPts, y = HomeWinMargin)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(title = "Scatter Plot of AwayPts vs HomeWinMargin")
## `geom_smooth()` using formula = 'y ~ x'

```

Scatter Plot of AwayPts vs HomeWinMargin



```

# Conclusion: Based on the correlation coefficients and the scatter plots, it appears that HomePts and ...  

# Follow-up question: Win margin is calculated using home and away points directly, is it trivial or re...  

# Answer: Using HomePts and AwayPts as predictors for HomeWinMargin may be considered trivial or redundant.  

# 3.1.1 - Avoiding Overfitting  

# Keep in mind that: All derby matches are played at a neutral venue with fan support split 50% for team ...  

# Derived features  

# - Ratio of HomePts to AwayPts  

# - Difference in ELO ratings between the two teams (using the ELO ratings calculated in Step 2)  

# - Fatigue index for both teams (number of games played leading up to the matchup)  

# - Historical performance of the teams in derby matches (e.g., average win margin in previous derby ma...  

# - Team strength indicators (e.g., average points scored and allowed throughout the season, ranking ba...
# - Possible ELO metrics, rolling averages of ELO ratings over the last 5-10 games, etc. To capture team ...

```

```

# - There needs to be guardrails put in place so the values for this are not incorrect. i.e. GameID 100
# - All/any metrics that are derived need to be calculated in a way that they are not using information
# - We can use the "lag" function to create features that are based on historical performance, ensuring
# - When we are creating derived features, we need to ensure that we are not introducing any data leaka
# - When using ELO derived features, each observation needs to reflect the correct index of the ELO rat

safe_derived_features = train %>%
  group_by(HomeTeam) %>%
  arrange(GameID) %>%
  mutate(
    RatioHomeAwayPts = HomePts / (AwayPts + 1), # Adding 1 to avoid division by zero
    ELO_Difference = ELO_RATINGS_v1$Rating[match(HomeTeam, ELO_RATINGS_v1$Team)] - ELO_RATINGS_v1$Rating,
    HomeFatigueIndex = cumsum(HomeTeam == lag(HomeTeam, default = first(HomeTeam))),
    AwayFatigueIndex = cumsum(AwayTeam == lag(AwayTeam, default = first(AwayTeam))),
    HistoricalWinMargin = ifelse(GameID > 1001, lag(cumsum(HomeWinMargin), default = 0) / (cumsum(HomeW
  ) %>%
  ungroup()

safe_derived_features

## # A tibble: 940 x 16
##   GameID Date   HomeConf HomeID HomeTeam     HomePts AwayConf AwayID AwayTeam
##   <dbl> <chr>   <chr>   <dbl> <chr>      <dbl> <chr>   <dbl> <chr>
## 1 1001 1/1/2025 Yellow  123 Maine       21 Red    85 New Hope
## 2 1002 1/1/2025 Green   53 Wyoming    77 Purple   70 Maryland
## 3 1003 1/1/2025 Purple  68 Idaho      123 Red    83 Jeffers-
## 4 1004 1/1/2025 Purple  74 Salt Lake C~  51 Crimson 22 Michigan
## 5 1005 1/2/2025 Gold   33 El Paso    41 Orange   57 Denver
## 6 1006 1/2/2025 Crimson 26 Phoenix   129 Yellow  117 Columbia
## 7 1007 1/2/2025 Green   46 Boston    125 Crimson 25 Oregon
## 8 1008 1/2/2025 Red    79 Des Moines 81 White   110 Raleigh
## 9 1009 1/2/2025 Gold   28 San Antonio 49 Gold    36 Massach-
## 10 1010 1/2/2025 Crimson 19 Georgia  81 White   115 Wiscons-
## # i 930 more rows
## # i 7 more variables: AwayPts <dbl>, HomeWinMargin <dbl>,
## #   RatioHomeAwayPts <dbl>, ELO_Difference <dbl>, HomeFatigueIndex <int>,
## #   AwayFatigueIndex <int>, HistoricalWinMargin <dbl>

```